Growing up in a compound for doctors due to my dad's job, and with friends studying medicine, I developed an interest in healthcare early on. The first time I worked on the healthtech intersection was after studying ML in my bachelor's. My mind went to how biases in healthcare literature would be cemented in the upcoming age of AI healthcare. I compiled my findings into a blog and messaged researchers working on such models at the time, proposing methods to check fairness. Received a few positive responses, but a lack of resources prevented any actual work. That's what motivated me to pursue a master's and be part of an institute again.

Satisfied with my ML knowledge base, I focused on applying it. I took a health innovations course covering the EU healthtech landscape and the end-to-end process of bringing innovations through it. A hospital project course investigating EHR interoperability challenges in Sweden, learning the gap between implemented and existing standards. A small consultancy project forecasting patient arrivals at an emergency department. Most recently, I've started a pilot with a hospital in Saudi Arabia on using AI to read electrophoresis results and speed up lab turnaround times. None are deep research projects individually, but the breadth of exposure means I understand the structures and common problems. A PhD coming from that awareness is more likely to produce something workers actually want to use, and that fits the system.

The pilot in Saudi Arabia is where I've been thinking most seriously about what good healthcare AI design actually requires. The graphs and scan data are heavily imbalanced, and simple accuracy or recall metrics would give false confidence in the system. What I find more interesting is talking with the physicians and mapping their regular workflow, then asking whether the AI addition would be a clear improvement. You also start seeing second-order problems that don't show up in papers. Radiologists, for instance, may find that once the straightforward cases are handled by AI, every remaining case they see is complex. My biggest motivation for picking a PhD position is knowing the research is designed with those realities in mind, that the right stakeholders are involved in decisions, and that the work has a realistic path to being used. That's the condition under which I know my time will pay off.

My master's thesis includes an annotated dataset of short-form video comments I built and a benchmark of classification approaches, including a two-stage multimodal pipeline for image-based comments. It required working directly with label imbalance, annotation disagreement and producing a custom architecture. I'm proud of that last part, because I noticed that it was easier to label comments as harmful or not based on the overall sentiment in the comment section, so why not include multiple passes in the prediction architecture?

My recent ML projects have mostly been in NLP or forecasting, so I'm not up to date on computer vision methods outside of my courses. Learning about the latest CNN architectures or YOLO is something I can read up on my own. The hard part is figuring out how and where we'll apply it once we start working with the data. The earliest I can start is August 2026.

Two of my references have attached letters. The third, Aria, was my supervisor for both semester projects in my final year at Uppsala, and I worked a bit with at Fyris AI. My degree is pending my thesis presentation on the 11th of June.